


Before and after: one prompt change, measured on 30 real cases

Dyuthi • 2026-09-08 • written in my own words, not a formal report

Average score went from **1.77** to **1.93** out of 3 across the same 30 cases, run through the real model both times. The number of comments that hit the post's actual distinctive detail nearly doubled, from 6 to 10, because the one thing I changed was telling the model to find that detail before writing instead of settling for the general topic. It didn't fix everything — the 3 cases where the honest answer is no comment at all (2 bereavement posts, 1 euphemistic layoff announcement) still scored 0 in both runs, because that's a different problem this change never touched.

Score distribution across all 30 cases (0 = worst, 3 = best)



	Baseline	After	Change
Average rubric score (0-3)	1.77	1.93	+0.16
Score of 3 (hit the specific detail)	6/30	10/30	+4
Score of 2 (generic but personalized)	14/30	11/30	-3
Score of 1 (no personalization at all)	7/30	6/30	-1
Score of 0 (blocked / wrong tone / SKIP case)	3/30	3/30	+0
Generic-phrase failures (gate's R4/TEMPLATE3 flag)	0/30	1/30	+1

What the one change actually was, and why

I expected the fix to be about banned phrases and the formulaic closing question — that's what the demo found two days ago. But I measured the baseline first, before touching anything, and that assumption turned out wrong: zero of the 30 baseline comments tripped a banned-phrase or 3-beat-template flag. The real problem, measured, was specificity — 14 comments passed every rule but talked about the general topic instead of the one distinctive detail, and 7 more showed no personalization at all.

So the one thing I changed was a new instruction in the prompt telling the model to find the single most distinctive detail — a number, an admission, a concrete consequence — before writing, and explicitly ruling out industry or job title as counting as “specific.” Nothing else moved: same banned-phrase list, same length limits, same everything else.

One after-run comment opened with “You make a solid point” — the exact banned-opener pattern from two days ago — and the gate correctly caught it. That's not this change failing; it's a different failure mode this change was never aimed at. Change one thing, measure it honestly, and don't credit it for problems it didn't touch.

Where everything is

docs/golden_set/BEFORE_AFTER_MEMO.md — the full write-up these numbers come from.

docs/golden_set/baseline_run.csv and **after_run.csv** — one row per case: the actual candidate text, the quality flags, the rubric score and why.

scripts/run_golden_set.py — rerun it and get the same numbers; the scorer is code, not a feeling.

src/outreach/copy.py — the one prompt change, in `_STYLE_RULES`.

Nothing here is committed or pushed yet. Say the word and I'll push it.

Still open, on purpose, not folded into this number: the 3 score-of-0 cases (2 condolence posts, 1 euphemistic layoff) are unchanged in both runs. Teaching the pipeline to recognize those and suppress the suggestion entirely is a bigger, separate change — it deserves its own before/after, not a line item borrowed from this one.