

Comment Quality Spec v1

Draft for review · comment action only · Dyuthi T G · 2026-08-25

“Good enough to publish” currently lives in Deep's head. This replaces it with rules that have a pass/fail test a machine can run — if a rule can't be written as a test, it isn't finished. **Automatic reject** if any of: sentence count $\neq$ 1–3 · length >280 chars · no specific reference to the source post · banned phrase · matches the formulaic closing-question shape · any emoji · >1 exclamation mark · unfilled placeholder / link / own-product mention (already enforced today).

Rules

Rule	Test	Why
R1 — 1–3 sentences	$1 \leq \text{sentence_count} \leq 3$	Not enforced anywhere today; prompt itself only asks for “one or two” (copy.py:139) — fix both.
R2 — $\leq$ 280 characters	$\text{len}(\text{comment}) \leq 280$	Not a guess: good-graded comments ran 161–258 chars, fail-graded ran 143–269. Length doesn't predict quality — this is a rambling backstop, not a filter.
R3 — specific to this post	≥ 1 non-generic 6+ letter word from the post (or a proper noun) appears in the comment	Fixes the audit's core bug: quality.py stoplists generic words for a headline match but not a post-text match, so 58/58 comments scored 100/100 personalization.
R4 — no banned phrase	case-insensitive substring match, list below	Seed list (3 phrases) extended from what actually recurred in the audit, cross-checked against the audit's own failure tags — not every frequent phrase made the cut.
R5 — no formulaic closing question	in a sentence ending ‘?’: pronoun near found/find/noticed/seen/tried, AND a topic noun (strategy/approach/tactic/framework/trend/metric) anywhere in it	43/58 audited comments (74%) ended in a version of this. ~15 literal wordings of one template exist — banning the shape beats whack-a-mole on strings.
R6 — zero emoji	no emoji characters present	Gate today allows up to 2; the rule says none.
R7 — ≤ 1 exclamation mark	$\text{comment.count}("!") \leq 1$	Gate today allows up to 2; the rule says one.

R3 calibration note: an early draft's generic-word list also included “revenue,” which wrongly rejected a good-graded comment about “net revenue retention” — the post's own subject. Test against data before adding to a stoplist, not instinct. **R5 known gap:** two audited comments (“have you noticed any shifts/changes in user engagement”) use the same template with a topic noun outside the current list — left as a v2 item rather than chased, to avoid memorizing this one dataset.

R4 banned-phrase list

thanks for sharing (seed) · great insights (seed) · spot on / you're spot on (seed) · you make a solid/great point · you raise a good/crucial point · you've nailed it / you nailed it · absolutely agree · that's a solid approach · i hope this message finds you well · quick question · circle back · touch base · pick your brain · synergy · game-changer

Deliberately left off despite recurring often: “it's interesting how,” “i've seen teams” — the human audit never tagged either as a violation, and one good-graded comment uses the first outright. Frequency isn't the bar; matching a named failure is.

Brand voice: axes, not four documents

One spec; voice varies along axes mapped to concrete rule/prompt changes — never a rewritten document per account. Replaces the free-text `ICPPProfile.instructions` field (src/targeting/models.py:85), which is how undocumented per-account drift happens today.

Axis	Range	Effect	Lake B2B	Ampliz	Champions Ranch	Deep's acct.
Formality	1–5	<3: contractions required, R1 favors 1–2 sent.	4	3	2	2
Warmth	1–5	>3: one non-R4 affirming opener allowed	2	3	4	3
Sentence len.	short/med	target words/sentence, R2 cap unchanged	medium	medium	short	short
Question close	never/sometimes	“never” extends R5 to any closing question	sometimes	sometimes	sometimes	never

Example values above are placeholders pending ten minutes with each account owner — not shipped as fact.

Validation — run for real, not asserted

scripts/validate_spec_v1.py grades comments against R1–R7 with no access to the human grade already in audit_log_scored.csv. **First run, before calibration: 6/10** agreement on a stratified 10-comment sample — it rejected 2 good comments on an over-tight length guess and missed a weak one on an under-built R5. Each miss traced to a specific bad assumption and fixed against data (see notes above), not intuition. **After calibration: 10/10 on the sample, 56/58 (97%) on the full audited set**, two named residual misses left open rather than hidden. Caveat: synthetic prospects through the real pipeline, not a live account — re-run once one is reachable before trusting this to gate production traffic.

What ships this week vs. what needs Hemang

This week, no new logic: extend _TIRED_PHRASES with the R4 list, fix the R3 stoplist gap in _personalization_signals(), fix the copy.py:139 prompt text. **Needs Hemang + a booked slot:** R1 and R5 are new logic inside quality.py that connect/message also depend on. **Needs account owners:** confirm real values for the brand-voice axis table.