

The 30-case golden set, and the rubric to score against it

Dyuthi • 2026-09-07 • written in my own words, not a formal report

The brief was: 30 real-shaped posts, the ideal comment for each written by hand, and a 0-to-3 rubric with anchors clear enough that someone else scores the same way I would. Slow on purpose. Here's what I shipped and where to find each piece.

1. Thirty posts, spanning what the bot actually sees

5 each of funding news, hiring posts, thought leadership, product launches, and personal milestones. Plus 3 layoffs and 2 condolence posts, on purpose, because you asked for the awkward ones and you're right that's where this breaks.

None of these are real people. I made them up. Two reasons: we still don't have live LinkedIn access to go scrape real posts with, and even if we did, I didn't want a permanent test file in the repo built out of real people's real grief and real layoffs. Each one is written to feel like a real post though — specific numbers, specific small details, not generic filler — because a vague fake post would let a vague fake comment pass, and that defeats the point.

2. A hand-written ideal comment for every one of the 30

I wrote what I'd actually post, not what sounds impressive. The pattern across almost all of them: respond to the one specific detail that makes this post different from the four others in its category, not the general topic. "Congrats on the raise" works on all five funding posts equally, which is exactly why it's wrong for all five.

3. The part that actually mattered: three cases where the right answer is silence

Two condolence posts and one layoff post announced in corporate-speak ("optimizing our workforce to align with strategic priorities") are marked SKIP in the notes column. For those three, I didn't write an ideal comment, because there isn't one — the correct output is no comment at all. A well-written comment on a bereavement post is not a slightly-wrong answer, it's the specific failure that costs you a person's trust permanently.

This is why the rubric's first question isn't about tone or specificity — it's "should this have gotten a comment at all." If the bot ever writes fluent, on-topic text for one of these three, that's a 0 no matter how good the sentence is. I'd rather find that out against 30 fake posts today than against a real one later.

4. The rubric: one number, four anchors

- **3** — exactly what I'd have posted myself, responds to the one distinctive detail
- **2** — good, read the post, right tone, but I'd tweak one phrase
- **1** — technically fine, but could've been posted on anyone's similar update unchanged
- **0** — wrong tone for the moment, a formulaic template shape, or commented on a SKIP case at all

No sub-scores for tone, specificity, length, and so on separately. One number. I tried a longer version first and it took four times as long to apply and didn't produce different scores than the short one — the anchors were doing the real work, not the extra dimensions.

5. How I'll know the rubric actually holds, not just feels right to me

I haven't skipped this, but I also can't finish it alone — it needs a second person. The plan, written down in the rubric file so it's ready to run: pick 3 posts (one easy, one medium, one hard), give Harshil only the source posts, no ideal comments, no notes, and have him write a comment for each blind. Then I score his three against the rubric myself.

If he agrees with my score and the reason I give for it, the rubric is anchored well enough to trust on real bot output. If he doesn't — I call his comment a 2, he thinks it's obviously a 3 — that's not him being wrong, it means one of the anchors is vague enough that two reasonable people read it differently, and I go fix that anchor with his exact case as the new example.

This is the one piece of today's brief I can point you to but haven't closed out, because it needs Harshil's actual time, not more of mine.

6. Where everything is

- **docs/golden_set/golden_set.csv** — all 30 cases: post, author, hand-written ideal comment, notes
- **docs/golden_set/RUBRIC.md** — the 0-to-3 rubric, the anchors, and the Harshil blind-test steps
- **scripts/build_golden_set.py** — regenerates the CSV from source; edit the posts here, not the CSV directly

Nothing here is committed or pushed yet. Say the word and I'll push it.

What I need from you: 15 minutes of Harshil's time for the blind test in section 5. That's the one thing that turns "I think this rubric is right" into "we checked, and it is."