

What ships today: the runbook, the voice guide, and the quality writeup

2026-09-16

The premise for today: I'm the bottleneck on this product as long as clearing the queue, tuning the voice, or rerunning the quality check only works if I'm the one doing it. Three documents, each aimed at someone who joins next month and has never met me.

1. The three asks, and what satisfies each one

Ask	New file	What it covers
Operator runbook	docs/OPERATOR_RUNBOOK.md	Clearing the queue (the four actions and what each does), the full flag glossary read straight from quality.py's rule table, what to do when a session expires, and a named escalation table.
Voice-tuning guide	docs/VOICE_TUNING_GUIDE.md	Adding/editing a brand profile (no deploy) versus changing a shared rule or the global banned-phrase list (code + deploy) -- the one distinction that determines who can make a given change without going through me.
Quality measurement method	docs/QUALITY_MEASUREMENT_METHOD.md	How the golden set, the rubric, and the runner fit together, the exact command to rerun it, and what the resulting number does and doesn't prove.

2. New pages versus pages that already existed

Two things already covered part of this ground, and the new docs point to them rather than repeat them: **docs/CONNECTING_AN_ACCOUNT.md** already has the account setup, warm-up stages, and a symptom/meaning/fix table for account problems.

backend/config/brand_voices/README.md already explains adding a brand profile field by field. What was actually missing, and is what the two new docs add: the queue-actions and flag glossary, the escalation table, and -- the real gap -- nothing anywhere explained the difference between a per-brand change (a marketer can do it, no deploy) and a shared-rule change (needs code and a deploy). That line is now explicit in **VOICE_TUNING_GUIDE.md** §2.

One real inconsistency found while writing this

Week3_And_Go_Live_Checklist.pdf's daily-check line says to run **validate_account.py --replace**. The flag actually implemented in the script is **--rotate** -- there is no --replace. OPERATOR_RUNBOOK.md §3 uses **--rotate** and calls out the discrepancy explicitly so whoever reads the older PDF first doesn't type a flag that doesn't exist mid-pilot.

3. The bar, and what's still owed before Friday

"Good" here isn't these documents existing -- it's handing the runbook to someone who has never used the tool and watching them clear five items with no help. That test needs a real second person at the keyboard, so it isn't done by writing the page; it's the next step, not a step I completed by producing this PDF.

- 1 Sit someone who's never opened Approvals down with only OPERATOR_RUNBOOK.md open.
- 2 Have them clear five real items -- approve, edit, skip, whatever the queue actually gives them.
- 3 Every question they ask out loud is logged verbatim. A question is a gap in the page, not a gap in them.
- 4 Fix each logged gap in the runbook before Friday, then consider it shipped -- not before.

Nothing above is committed or pushed yet -- same as always, tell me and I'll do it.