

# Month 1 review: what I built, what it's worth, what's next

2026-09-18 • 15-minute review for Deep

## 1. The 15-minute demo, in order

- 1 The problem (2 min).** The bot had been stuck since March: a review queue nobody trusted, quality nobody had measured, and no path from “it runs” to “someone would actually let it touch a real account.”
- 2 What I built (7 min).** A real review queue wired to real data with keyboard actions and undo; a quality gate with a named rule table, shown the same way to the reviewer and in the spec; brand voice as a per-account config file, no deploy needed; guardrail flags visible in the queue instead of hidden in logs; an operator runbook, a voice-tuning guide, and a documented, rerunnable quality measurement method.
- 3 The numbers (3 min).** Section 2 below.
- 4 What's still open (3 min).** Section 3 below -- said plainly, not glossed over.

Rehearsed once, out loud, alone, before this goes in front of Deep -- per the brief.

## 2. The before/after quality numbers

Method: `backend/scripts/run_golden_set.py` runs the real copywriter against all 30 golden-set cases (5 each of funding / hiring / thought-leadership / product-launch / personal-milestone, plus 3 layoffs and 2 condolences that should get no comment at all), through a real OpenRouter call, and grades every result with a deterministic 0-3 scorer. Same code, same 30 cases, same scorer, both times -- that's what makes the comparison honest.

|                                                 | Baseline | After one prompt change | Change |
|-------------------------------------------------|----------|-------------------------|--------|
| Average rubric score (0-3)                      | 1.77     | 1.93                    | +0.16  |
| Hit the post's specific hook (score 3)          | 6 / 30   | 10 / 30                 | +4     |
| Generic but personalized (score 2)              | 14 / 30  | 11 / 30                 | -3     |
| No personalization at all (score 1)             | 7 / 30   | 6 / 30                  | -1     |
| Blocked / wrong tone / SKIP-case miss (score 0) | 3 / 30   | 3 / 30                  | 0      |

The one change tested was a prompt instruction telling the model to find the single most distinctive detail in the post and build the comment around it. It measurably helped specificity (the

“3” column nearly doubled) and, by design, didn't touch the three score-of-0 cases -- the bot still has no concept of “don't comment on this post at all,” which is a separate, bigger change. Verified again today as part of this review, rerunning the same script live: **2.00/3** on a fresh run -- same range, real variance from a real model call, not a fabricated repeat of the old number.

Caveat that stays attached to every one of these numbers: this is golden-set data -- realistic but synthetic, not a live account. It measures “did this change help in principle,” not a substitute for reviewing real queue output.

### 3. Is the bot live?

**Not yet. Named blocker, not a vague “almost there.”**

The Sep 15 brief called for a real person approving 20 real items on a real disposable account, with Deep watching and Harshil's written sign-off on the blast-radius plan on record first. That session hasn't run. Three prerequisites are still not confirmed green: a real **ENCRYPTION\_KEY** (Hemang), Redis actually standing up with /healthz reporting ok (Deep), and Harshil's written sign-off itself. Full detail in **Sep15\_Live\_Dry\_Run\_Status.pdf**.

**Recommendation:** once those three items are confirmed green -- which is infra work, not product work, and doesn't depend on me -- the live session can run within the same week. Proposing **Sep 22** as the target date for the real 20-item session, with the actual go/no-go to follow once that log is filled in, not before.

### 4. Retro: what worked, what I'd redo

#### What worked

- Measuring the baseline before guessing at a fix. I assumed the dominant problem was still yesterday's banned-phrase issue; the baseline measurement proved that assumption wrong and pointed at specificity instead. Guessing first would have spent a day fixing the wrong thing.
- Changing one thing at a time and measuring it. The before/after number is trustworthy specifically because nothing else moved between the two runs.
- Reusing what already existed instead of inventing parallel systems -- the existing account-activity audit ledger for approvals, the existing BREAK\_LIST.md shape for issues, the existing ship-note format for every PDF. Fewer formats for the next person to learn.

#### What I'd redo

- I let the Sep 15 live-session prerequisites (ENCRYPTION\_KEY, Redis) sit unconfirmed for three days before naming them plainly in a status doc. Naming a blocker the day it's discovered, not the day someone asks about it, is the actual habit worth keeping.

- The Sep 16 repo split (frontend/ and backend/) shipped without an audit of which scripts compute paths relative to their own location -- one (run\_golden\_set.py) broke silently and only surfaced today because I happened to run the exact command the docs tell people to run. A structural change like that deserved a five-minute grep for Path(\_\_file\_\_) across scripts/ before calling it done.
- Still only one reviewer on the whole project. That was fine at demo scale; it's the first real limit Month 2 hits if account count goes up before a second approver is trained.

## 5. Month 2 proposal

**One recommendation:** don't widen to more accounts until the Sep 15 live session has actually run and passed go/no-go. Once it has, add a second trained reviewer before adding a second account -- the single-reviewer bottleneck in **docs/REVIEW\_SLA.md** §2 is the more urgent scaling limit than account count, and it's cheaper to fix first. In parallel, start the ChampESS front-end track (per the program's own Month 2 window, 21 Sep to 16 Oct) since it doesn't compete for the same reviewer time.

**What I'd stop doing:** building new guardrail-flag types ahead of demand. Every flag in the current rule table traces back to a real thing the audit or the golden set found -- that discipline is worth keeping instead of adding rules that feel plausible but aren't backed by a measured failure.

The Ranch and Montana idea, parked in August, is on the table for Month 2 per the program's own plan -- I don't have the specifics of that idea written down anywhere in this repo or in what I've worked from this month, so bringing it back out needs a real conversation with Deep, not a guess from me about what it was.

Nothing above is committed or pushed yet -- same as always, tell me and I'll do it.