

Task Audit — Brand Voice / Social Copy Auditor

Dyuthi T G • LinkedIn Automation Platform • 2026-08-17 to 2026-08-25

This is a rollup of two completed audits: whether the bot's LinkedIn comments actually read as human, and whether the rules that are supposed to guarantee that are actually enforced by code rather than only requested in a prompt. Every finding below cites a file, a line, or a real generated comment — nothing here rests on opinion alone.

1. Comment Quality Audit (2026-08-20 – 2026-08-21)

Objective

Go from zero visibility into the bot's real LinkedIn comments to a scored, ranked failure taxonomy that Harshil and Deep can act on.

What I did

- No audit-log CSV export existed and no comment history was reachable in a live account.
- Built **scripts/audit_run.py**: generates a real 58-comment batch through the actual production pipeline (targeting → copywriter → quality gate) against 58 synthetic prospects.
- Diagnosed why the LLM wasn't firing (comments were silently falling back to templates): an import in the harness was wiping **OPENROUTER_API_KEY** at module load. Fixed by inlining the recording transport.
- Found the configured model (**anthropic/claude-3.5-sonnet**) is not a valid OpenRouter endpoint; switched to **openai/gpt-4o-mini** to get real generations.
- Exported **audit_log.csv**, built **scripts/score_audit.py** with a self-designed structural failure taxonomy, and hand-scored all 58 comments, including the good ones.
- Traced the root cause of why the product's own quality gate missed all of this: a specific bug in **src/outreach/quality.py**.

Headline finding

The bot writes one comment: **validate** → **restate the thesis** → “**have you found any specific strategies?**” 43 of 58 comments (74%) end in that exact question shape; 25 of 58 run the full three-beat skeleton. Individually each comment reads fine — stamped on 58 different people from one account, it is an unmistakable bot fingerprint. The deterministic quality gate scored **all 58 of them 100/100**.

The specific code bug (not a summary — a named defect)

src/outreach/quality.py, **_personalization_signals()**, lines **300–338**. The headline-match check stoplists generic business vocabulary before crediting a match (lines 324–327, via **_GENERIC_HEADLINE_WORDS**). The post-text match three lines below it (330–336) applies **no stoplist at all**: any shared 6+ letter word between the comment and the source post is credited as personalization. Every audited comment was about activation, retention, or onboarding because the post was — so this check passed automatically almost every time, including on the templated “you nailed it...have you found any specific strategies” comments.

Verdict split (58 comments)

Good	Weak	Fail
6 (10%)	41 (71%)	11 (19%)

Ranked failure taxonomy (frequency × severity)

Rank	Failure	n	Freq×Sev
1	formulaic_strategy_question	43	2425
2	validation_opener	32	2010
3	template_3beat	25	1650
4	no_specific_detail	20	1384
5	sycophancy	10	610
6	exclamation_hype	2	100

Worst comment, verbatim (score 5/100)

“You’re spot on about the ICP issue. It’s interesting how a narrow focus can actually improve both activation and retention. Have you found any specific strategies that work well for refining ICP in a B2B context?”

What I recommended (Section 1)

- Stoplist the post-text personalization check the same way the headline check already is (quality.py:330–336) — a one-line fix that closes the exact hole above.
- Add an account-level repetition check: a cheap structure-overlap check across an account's pending suggestions, run before send, since `check_copy()` is stateless per call and cannot see the pattern across people.
- Extend `_TIRED_PHRASES` to the opener/closer stems found here (“you make a great point,” “have you found any specific strategies,” “spot on”) — the pattern-list mechanism already exists, it just isn't pointed at these phrases.

Deliverables shipped: `audit_log.csv`, `audit_log_scored.csv`, `audit_report.md`, `BOT_OUTPUT_AUDIT_V1.md`, `RAMP_WEEK_REPORT.md`, `scripts/audit_run.py`, `scripts/score_audit.py`.

2. Rules-vs-Reality Enforcement Audit (2026-08-24 – 2026-08-25)

Objective

The documented comment rules are good rules — 1–3 short sentences, a hard cap around three lines, a banned-phrase list, must reference something specific. Nobody had checked whether the machine actually follows them, as opposed to merely being asked to.

Where the rules actually live in code

The path named in the original brief, `services/comment_generator.py`, does not exist in this repository — flagged for Harshil to confirm whether that's a stale reference. The real split is:

- **Prompt** — `src/outreach/copy.py`: `_STYLE_RULES` (lines 37–52) and `_task_for()` (lines 119–144). This is what the model is asked to do; nothing here is enforced.
- **Review pass** — `src/outreach/quality.py`: `check_copy()` (lines 151–297) and `_personalization_signals()` (lines 300–338). This is the only place a rule becomes real.

Rules-versus-reality gap table

Rule	Enforced?	Evidence
1–3 short sentences	No	No sentence-splitting/counting logic exists in <code>check_copy()</code> (quality.py:151-297); prompt itself says “one or two” (copy.py:139), not one to three.
Hard cap ~3 lines	Proxy only, ~2x too loose	<code>COMMENT_MAX</code> = 400 chars (quality.py:34) ≈ 6 lines at 65 chars/line. No code counts lines or <code>\n</code> .
Banned phrase list	Mechanism yes, coverage no	<code>_TIRED_PHRASES</code> (quality.py:38-70) is checked (228-231) but has no entry for “spot on,” “nailed it,” “have you found any specific strategies,” or “great post.”
Must reference something specific	Half-enforced	Stoplist applied to headline match (quality.py:323-327) but not to post-text match (330-336) — the exact bug in Section 1.
No flattery / no product mentions	No	No such check exists anywhere in <code>check_copy()</code> ; matches the sycophancy tag on 10/58 audited comments.

Sorted by what it takes to fix

Prompt problems — fixable this week, no new logic:

- Extend `_TIRED_PHRASES` / `_GENERIC_OPENERS` with the phrases actually found in output. The list mechanism already runs; this is data entry.
- Fix the prompt text mismatch at copy.py:139 (“one or two” → “one to three” sentences).

Missing-code problems — need Hemang and a booked slot:

- No sentence-count check exists at all — needs new parsing logic in `check_copy()`.
- The char-cap-as-line-cap proxy is shared across connect/message/comment (`_limit_for()`, quality.py:143-148) — tightening it isn't isolated to comments.
- The post-text stoplist fix is small, but src/outreach/quality.py isn't mine to edit unilaterally — ownership decision already logged as an open ask.
- A real flattery/own-product detector doesn't exist in any form and needs design (the gate has no concept today of the account's own product/company to check against).

Deliverable shipped: `COMMENT_RULES_ENFORCEMENT_GAP.md`.

3. Forward-looking: proposed differentiation features (not yet scoped)

Raised for roadmap discussion, not built or estimated yet. Each targets the same root problem both audits above found: not getting flagged is not the same as being recognizably human.

- **Per-account voice fingerprinting** — derive style from that account's own real posts instead of one global prompt for every user.
- **Adversarial “would this get me caught” pass** — a second model judges the draft as the recipient would, catching shape problems regex-based gates miss.
- **Deliberate, occasional disagreement** — counters the uniformly-agreeable pattern this audit found in nearly every comment.

- **Relationship continuity across interactions** — reference prior engagement with the same person instead of treating every comment as stateless.
- **Outcome-linked style learning** — reweight generation based on which comment shapes actually get replies, not just which ones pass the gate.

4. Open asks / decisions needed from you

- No real audit-log CSV export exists yet, and no comment history is reachable in a live account — every finding above is synthetic-prospect data through the real production pipeline. Recommend re-running against a live account once one is reachable to confirm the pattern holds.
- Is the quality-gate stoplist fix (Section 1) mine to ship, or does it go to whoever owns src/outreach/quality.py?
- The sentence-count and char-cap fixes (Section 2) touch logic shared with connect/message — needs Hemang's sign-off on scope before anyone starts.
- Confirm with Harshil whether services/comment_generator.py refers to a different/older codebase, since that path doesn't exist here.

5. All deliverables shipped this period

File	What it is
DAILY_WORK_LOG.md	Day-by-day log of the comment quality audit build
RAMP_WEEK_REPORT.md	Week 1 summary + 15-minute demo script
BOT_OUTPUT_AUDIT_V1.md	Primary audit: ranked taxonomy, verdict split, verbatim worst/best, quality-gate defect
audit_log.csv / audit_log_scored.csv	Raw and scored data underlying the audit (58 comments)
audit_report.md	Raw scorer output
scripts/audit_run.py	Re-runnable: generates comments through the real pipeline
scripts/score_audit.py	Re-runnable: applies the structural failure taxonomy
COMMENT_RULES_ENFORCEMENT_GAP.md	Rules-vs-reality gap table, code locations, prompt-vs-missing-code sort
TASK_AUDIT_2026-08-25.pdf	This document

Provenance note: Section 1 and 2 findings are from the real production pipeline (targeting → copywriter → quality gate) run against synthetic prospects, since no live-account audit log was reachable. The taxonomy, the quality-gate defect, and the enforcement gaps are real regardless of data source; do not present the specific phrasing distribution as live-account data.