

# Task update

Dyuthi — brand voice / social copy audit — Aug 17 to 25

---

Here's where things stand on the comment audit work. Two pieces: the quality audit I ran last week on what the bot is actually posting, and a follow-up this week checking whether the comment rules we already agree on are enforced in code, or just requested in the prompt and hoped for. Writing it up now so you've got the specifics before Friday rather than the summary version.

## What's actually coming out of the bot

There was no audit-log export to pull from and nothing reachable in a live account, so I built a script (scripts/audit\_run.py) that pushes 58 synthetic prospects through the real pipeline — targeting, then the copywriter, then the quality gate — and logs everything that comes out the other side. First pass produced identical template copy for all 58, which turned out to be my own fault: an import in the test harness was wiping the OpenRouter key before the model call ran. Fixed that. Also found that the model we have configured, claude-3.5-sonnet, isn't a valid OpenRouter slug — this run used gpt-4o-mini instead, so treat the exact wording below as illustrative, not final.

Once real generations were flowing I scored all 58 by hand against a taxonomy I built for this (scripts/score\_audit.py). Split came out 6 good, 41 weak, 11 fail. The thing that actually matters: 43 of the 58 — 74% — end with some version of the same question, “have you found any specific strategies for X?”, usually preceded by a “you nailed it” opener. That isn't 58 individual weak comments. It's one comment reused on 58 different people. And the quality gate scored every one of them 100 out of 100.

Worth being specific about why the gate missed it, because it's a fixable bug, not a fuzzy calibration problem. The personalization check stoplists generic words like “growth” or “strategy” before crediting a match against someone's headline (quality.py:324–327) — that part is genuinely well built. The equivalent check against the post text, three lines further down (quality.py:330–336), has no stoplist at all. Since every comment in this batch is about activation or retention because the source post was, that check auto-passes almost every time. That's the actual reason this needed a human doing it by hand.

| Rank | Failure                    | n  | freq × sev |
|------|----------------------------|----|------------|
| 1    | formulaic closing question | 43 | 2425       |
| 2    | validation opener          | 32 | 2010       |
| 3    | full 3-beat template       | 25 | 1650       |
| 4    | no specific detail         | 20 | 1384       |
| 5    | sycophancy                 | 10 | 610        |
| 6    | exclamation hype           | 2  | 100        |

Worst one, word for word:

*“You're spot on about the ICP issue. It's interesting how a narrow focus can actually improve both activation and retention. Have you found any specific strategies that work well for refining ICP in a B2B context?” — scored 5/100 by hand, 100/100 by the gate.*

What I'd fix first: stoplist the post-text check the same way the headline one already is — one line, in `quality.py`. Add something that looks across an account's pending comments for repeated structure, since right now each comment is scored alone and can never see the pattern. And add the specific phrases this run turned up (“spot on,” “have you found any specific strategies”) to the banned-phrase list that already exists and already runs — it just isn't pointed at these.

Everything above lives in `BOT_OUTPUT_AUDIT_V1.md`, with the full data in `audit_log_scored.csv`.

## Are the rules we have actually enforced

This week's follow-up was narrower: take the four rules everyone already agrees are good — one to three sentences, roughly a three-line cap, a banned-phrase list, has to reference something specific — and check whether the code actually holds that line, rather than just asking the model nicely and hoping.

One note before the table: the brief pointed at `services/comment_generator.py`, which doesn't exist anywhere in this repo. Flagging it for Harshil — either the brief is describing an older layout or he means a different file, worth two minutes before Friday. What does exist is a clean split: the prompt lives in `src/outreach/copy.py`, the actual enforcement lives in `src/outreach/quality.py`. Every row below is one rule, whether it's enforced, and the file and line proving it either way.

| Rule                              | Enforced?             | Evidence                                                                                                                                                                                   |
|-----------------------------------|-----------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1–3 short sentences               | No                    | No sentence-count logic anywhere in <code>check_copy()</code> ( <code>quality.py</code> :151–297); the prompt itself asks for “one or two,” not one to three ( <code>copy.py</code> :139). |
| ~3-line cap                       | Proxy only, ~2x loose | <code>COMMENT_MAX = 400</code> characters ( <code>quality.py</code> :34) works out to roughly 6 lines at 65 chars/line. Nothing counts lines or newlines.                                  |
| Banned phrase list                | Runs, but incomplete  | <code>_TIRED_PHRASES</code> ( <code>quality.py</code> :38–70) is actually checked (228–231) but has no entry for “spot on,” “nailed it,” or “great post.”                                  |
| Must reference something specific | Half enforced         | Stoplist applied to the headline match ( <code>quality.py</code> :323–327), not to the post-text match (330–336) — the bug from the section above.                                         |
| No flattery / no product mentions | No                    | No check for either exists in <code>check_copy()</code> at all.                                                                                                                            |

Two of these I can just fix myself this week — they're list edits, not new logic: add the missing banned phrases, and fix the one-word mismatch in the prompt where it says “one or two” sentences instead of “one to three.” The other three need actual new code in a file that connect and message also depend on, so those go to Hemang rather than me shipping them solo. Full table and reasoning is in `COMMENT_RULES_ENFORCEMENT_GAP.md`.

## Ideas worth floating, not built

A few directions came up while doing this that might be worth a longer conversation later. Not scoped, not estimated, just noting them here so they don't get lost: writing in a distinct voice per account instead of one global style for everyone; a second pass where a model judges the draft the way the recipient actually would, since that's exactly the kind of thing a structural gate keeps missing; letting the bot push back occasionally instead of only ever agreeing, since “always validates” is itself part of the fingerprint the audit found; and tracking which comments actually get replies so the system can learn from outcomes instead of only from the gate.

## **Where I need a call from you**

- Everything above is synthetic prospects through the real pipeline, not a live account — want to re-run once a live account is reachable to confirm the pattern holds.
- Is the personalization stoplist fix mine to ship, or does it route through whoever owns quality.py?
- The rule-enforcement fixes that touch shared code need Hemang's time booked, not just a heads-up.
- Need two minutes with Harshil on the comment\_generator.py path mismatch.

## **Everything that's on disk**

DAILY\_WORK\_LOG.md and RAMP\_WEEK\_REPORT.md (last week's log and summary),  
BOT\_OUTPUT\_AUDIT\_V1.md (the comment audit, full writeup), audit\_log.csv / audit\_log\_scored.csv /  
audit\_report.md (the underlying data), scripts/audit\_run.py and scripts/score\_audit.py (both re-runnable), and  
COMMENT\_RULES\_ENFORCEMENT\_GAP.md (this week's enforcement audit).

Happy to walk through any of this live if that's easier than reading it.