

Week 3: the numbers, the checklist, and the first live account

Dyuthi • 2026-09-11

1. The numbers first, because I've earned that this week

The one prompt change I tested this week moved the real quality score from **1.77** to **1.93** out of 3, on the same 30 real cases, run through the real model both times, not a guess. The comments that actually land on the post's one distinctive detail nearly doubled, 6 to 10. Full numbers are in **Before_After_Quality_Memo.pdf** already on your desktop — I'm not repeating the chart here, just leading with the headline because it's real and it's this week's.

Second number, less flattering but just as real: every single boot of this app on Railway was crashing before it served one request — the migration files just weren't in the image. That's now three small, boring fixes (copy the migration files in, add the Postgres driver, stop crashing when an old campaign has no organization) and it boots clean. I'm including it because "it's up" is a before/after number too, and it was false all week until Wednesday.

Third: brand voice profiles (below) have 23 real tests passing, and the people-and-post finder from last week has 8. I ran both again this morning before writing this, not just trusting my memory of the last time they passed.

2. What shipped this week

Account & run-status board. You can now see, per account, when it last ran, what broke, how much of today's sending limit is used, and whether it's in quiet hours — with a loud banner instead of a status dot you have to guess about.

Targeting discovery. Give it a profile link and it finds the person and their most recent post on its own. Before this, someone had to paste the post text in by hand or the queue came back empty with no explanation.

Guardrail flags in the queue. Every held-back comment now says exactly which rule it broke, in plain words, and you can filter the queue by reason instead of opening each one.

Brand voice as a config file. A marketer can now give a brand its own voice — how formal, how many sentences, which words to avoid, real example comments — by saving a file, no engineer and no deploy involved. Two real ones ship today: one formal (Lake B2B), one casual (Loopwork), proving the same six settings stretch to cover both. Built and tested (23 tests, all passing) but not committed yet — same as always, I'm waiting on your go-ahead before it goes in.

Two review-screen states that used to look identical. "Nothing generated yet" (a brand-new account, nothing has run) and "You're caught up" (plenty has run and it's all

decided) used to both show the exact same blank screen. They're now told apart. Built today and unit-tested; I have **not** clicked through it live in the browser yet, so I'm flagging that honestly rather than calling it proven.

3. The demo, in the order I'd run it

- 1 Open the account board first — it's the thing people were squinting at a status dot for last week.
- 2 On Targeting, paste in a profile link live and let it find the person and their latest post in front of you, not a screenshot of it working earlier.
- 3 In the queue, filter by one guardrail reason and show a held-back comment next to why it was held back.
- 4 Pick one target, generate a comment with no brand voice set, then set a brand voice on that same profile and generate again — same person, same post, two visibly different comments.
- 5 Close on the golden-set number from Section 1. It's the one thing in this whole demo that's a measured before/after and not a screenshot of a feature.

I'll reuse **Account_Health_Live_Walkthrough.webm** for step 1 since nothing changed there, and record fresh for steps 2 through 4 since brand voice and the frontend side of discovery are new since the last recording.

4. Go-live checklist — one name per line, on purpose

Every line below has one person's name, not a team's, because "Ops" can't tell me it's done and a name can. Some of these are mine and some genuinely aren't — I'm keeping the ones that aren't in, because a checklist that only lists my own work isn't a checklist, it's a diary.

Owner	What
Dyuthi	Create one brand-new, disposable LinkedIn account — not anyone's real profile.
Hemang	Set a real ENCRYPTION_KEY in the environment and confirm /healthz reports credentials_encryption: ok.
Deep	Stand up the Redis instance for this environment and confirm /healthz reports redis: ok. Sending stays refused until this is true — that's by design, not a bug to route around.
Dyuthi	Run scripts/validate_account.py against that account's cookies and get all four preflight checks green (whoami, profile, inbox, activity) before anything writes.
Dyuthi	Enroll the account in the warm-up program from day zero. No skipping stages because the profile looks established — that's the exact thing that gets accounts restricted.
Harshil	Read and sign off on the blast-radius plan in Section 6 before the first real send happens.

Owner	What
Dyuthi	Send the first single comment, approved by a human, and watch it happen live end to end — not queued and checked later.
Dyuthi	Check that account's session is still valid every morning of the pilot. Re-run <code>validate_account.py --replace</code> the moment it isn't, before anything else runs that day.
Kethan	Sit through at least one live approval session and say, as someone who didn't build it, whether the queue makes sense.
Harshil & Deep	Go/no-go: decide whether to move past one burner account, based on what actually happened, not on the plan.

If any line above can't get a plain yes or no from the name next to it, it isn't done — it's a wish.

5. The risks, stated plainly

Nothing below is hypothetical. Each one is either something the code already assumes will happen, or something the golden-set measurement already caught happening.

Account bans and restriction

The warm-up pacing, the daily caps, and the acceptance-rate safety net in the code all exist because LinkedIn restricts accounts that move too fast or too mechanically — that's not a maybe, it's the reason those three systems were built before a single live send. Even with all of them on, the honest position is that a burner account can still get flagged, and that's why it's a burner: I'm treating the first account as something we're allowed to lose, not something we're betting on keeping. If it gets restricted during the pilot, that's the plan working as intended, not the plan failing.

Terms-of-service exposure

This whole tool talks to LinkedIn the way their own app does, not through anything LinkedIn issued us permission to use. There's no version of this that's inside their terms — that's true today, was true in week 1, and doesn't get fixed by more engineering. It's a business risk to accept with eyes open, not a bug for me to close.

Session expiry

The login is a browser cookie, not a password, and it can stop working with no warning — a LinkedIn security check, a logout somewhere else, or just time. Today, nothing watches for that on its own; a human has to notice and re-run **`scripts/validate_account.py --replace`**. That's why checking it is a daily line item in the checklist above and not an assumption.

A wrong comment on a sensitive post

Measured, not guessed: in this week's 30-case test, the 3 cases where the right answer was no comment at all — two bereavement posts, one layoff announcement written in

euphemisms — scored 0 in both the baseline and the improved run. The bot has no concept yet of "don't comment here," and this week's prompt change was never going to fix that because it's a different problem. Until that's built, a human reviewing every single comment before it sends isn't a nice-to-have step in the blast-radius plan below — it's the only thing standing between the bot and exactly this mistake.

6. The first live account, and how small I'm keeping the blast radius

My recommendation: one brand-new, disposable LinkedIn account. Not Kethan's real profile, not mine, not anyone's real one. It's the account referenced as owner in Section 4 — created for this, and nobody minds losing it.

The blast radius, exactly:

- One account. Not a batch, not a small group — one.
- One send at a time. Nothing goes out in bulk during the pilot.
- A human approves every single item before it sends. No exceptions for "obviously fine" ones — that judgment call is exactly what Section 5's last risk says isn't safe to automate yet.
- The autonomous scheduler stays off. It already ships off by default; the pilot doesn't turn it on.
- Daily caps, quiet hours, and the warm-up program's pace all stay on from day zero, at their normal settings — nothing gets loosened to move faster.

I'm proposing this as the starting point, not the permanent shape of it. The right time to widen it — more accounts, batches instead of one at a time, less than 100% human review — is after the Sep 18 go/no-go, and only based on what actually happened Sep 16–17, not on how confident the plan reads today. I'd rather earn a bigger blast radius with a clean pilot than assume one and find out I was wrong on someone else's timeline.

Nothing new in Section 2 is committed or pushed yet — same as always, tell me and I'll do it.