

Today's work — the whole thing running, start to finish

Dyuthi • 2026-09-04 • written in my own words, not a formal report

Three things to ship today: run the whole loop on a staging account, record a short walkthrough, and write down everything that broke — plus the week 2 report. Here's what actually happened, starting with the part you'd want to poke at first.

First, the bit I'd rather tell you than have you find

The comment in the demo really does go all the way — the system finds the person, finds what they posted, writes the reply, checks it, I read it and approve it, and it gets sent. The sending part is a real request, built by the real code, with the real login headers.

But that request goes to a little pretend LinkedIn I wrote, running on my own machine. **Not to the actual LinkedIn.** We still have no LinkedIn account we're allowed to test against, and nothing in this whole system has ever touched the real thing. That hasn't changed since week 1, and it's the single most important sentence in this document.

I could have hidden that behind a fake “it worked!” message. I'd rather the pretend LinkedIn be obvious, because it means the parts we CAN prove are genuinely proven: the request really gets built, really gets sent, and the answer really gets read back. That's where the breakages hide, and a pretend shortcut would have proved none of it.

1. The loop runs, end to end, in about fifteen seconds

One command does the lot: find the people, find their posts, write the comments, flag them, show me the queue, take my approval, and send. There's a “practice mode” that does everything except the sending — I ran that first, which is how most of today's problems got found before anything was recorded.

2. The biggest thing I found: a whole step was missing

The system can't comment on a post unless it knows what the post says. Fair enough. But **nothing in the system was ever going and getting the post.** It only worked if a human had already typed the post into the spreadsheet by hand.

It looked like a finished feature because every test and every script we'd written happened to supply the post itself. You only notice when you try to run the real thing with nothing but a list of profile links — which is what a real list looks like. So I built that step today. Give it a column of links, it goes and finds out who those people are and what they recently posted. It only reads, it never posts anything, and it won't offer up a post we've already commented on.

3. Our own rulebook had never been typed into the code

Back on 25 August I wrote down a list of banned phrases and marked it “easy, do this week, no new code needed.” It never got done. So for ten days the written rule and the actual check were two different documents.

The run made that obvious in about a second: **three of the first five comments opened with “You make a solid point...”**, which is right there on my banned list, and the checker said nothing at all. Now it's wired in, plus a test that fails if any of those phrases ever quietly stops being checked.

4. My own test rig lied to me, so I made it argue with itself

The very first practice run said “no problems found” — while approving a comment that started by flattering the person and ended with the exact copy-paste question my August audit is named after. Scored 100 out of 100. Green light.

A test that can wave that through isn't testing anything. So now, after the comment is approved, the run grades it a second time using the completely separate scoring I wrote in August — which knows nothing about the main checker. **When the two disagree, the run records it as a problem.** They've disagreed on every single run since, which is exactly the point.

5. Everything that broke — twelve things, written down

You said the break list is worth more than the demo, and honestly it was. Six of them I fixed today with tests; five are written up with what I'd do about them and who needs to decide; one isn't really a bug at all. A few of the good ones:

- A rule that catches “*have you noticed any...*” didn't catch “*did you notice any...*” — same sentence, different tense. Fixed.
- “NRR?” got flagged as SHOUTING but plain “NRR” didn't — a question mark was tipping a three-letter abbreviation over a length limit. Fixed, and I added the abbreviations our own customers actually use.
- The bot wrote “At Champions Ranch, we've seen...” into a stranger's comment thread — advertising ourselves, which my own rules ban — and nothing stopped it. Not fixed: the checker doesn't currently know our own company name. Written up.
- If you skip a comment and then refresh the page within five seconds, **the skip silently doesn't happen** and the comment comes back. I only found this because it happened to me while recording. Written up with a fix.

- Importing a plain list of profile links gives you an empty queue and no explanation, because the system scores people before it knows anything about them. Written up.

6. The thing I most want you to know

The moment I banned “You make a solid point”, the very next run came back with **four out of five comments opening with “It’s interesting how…”** instead. Same job, same canned feeling, not on any list.

That’s my August argument happening in front of me inside one working day: banning exact phrases is whack-a-mole, and banning the *shape* of the sentence is the thing that actually holds. It’s why next week I’d rather widen the shape rule than keep adding words to a list.

7. The video — 2 minutes 19

Real screen, real comments the model wrote this morning, real posts the system went and found. It shows me skipping one that’s flagged, typing why, approving a clean one, and then the one from earlier that went all the way. The words on screen are what I’m thinking, not what the buttons do. It ends on the pretend-LinkedIn caveat, on purpose.

Two honest notes. It’s captions, not my voice — I’ve written the script out separately if you want it read aloud. And it’s a .webm file, not .mp4, because the video tool on this machine can’t make mp4s; it plays fine in any browser or VLC.

I had to re-record it twice. First take, the script approved a comment while my own caption said “this one’s clean” — it wasn’t, it had two flags on it. A demo whose narration argues with its own screen is worse than no demo, so I made the script read the real queue and pick a genuinely bad one and a genuinely good one before it says a word.

8. Where everything is

- **Week2_Demo_Walkthrough.webm** — the video
- **BREAK_LIST.md** — all twelve, with evidence for each
- **WEEK2_REPORT.md** — the week 2 report, gaps first
- **docs/DEMO_NARRATION.md** — the script, timed, if you want voice
- **scripts/staging_run.py** — the whole loop, re-runnable, with practice mode
- **staging_run/** — what the last run did, and every LinkedIn call it made

277 tests pass, including 13 new ones today. Nothing is committed or pushed yet — it’s all sitting here, tested and running. Say the word and I’ll push it.

And the one thing I need from you, not from more of my time: a burner LinkedIn account we're allowed to break, so next week that pretend LinkedIn can become the real one.